

[Feature flags and AI Configs](#) > [AI Configs](#)

Monitor AI Configs

Overview

This topic explains how to monitor the performance of your AI Configs. Performance metrics for AI Configs are available in the LaunchDarkly user interface if you [track AI metrics in your SDK](#).

Collect data for AI Config variations

Data appears on the **Monitoring** tab for an AI Config when you record metrics from your AI model generation. For each AI SDK, the function to record metrics takes a completion from your AI model generation, so you can make the call to your AI model provider and record metrics from model generation in one step. You can record duration, token usage, generation success and error, time to first token, output satisfaction, and more.

To learn how, read [Tracking AI metrics](#). For a detailed example, refer [Step 4 in the Quickstart for AI Configs](#).

☆ You can also monitor variation rollouts

You can use a [guarded rollout](#) when you are releasing new variations to your customers. The **Monitoring** tab shows a guarded rollout's progress and metric results. LaunchDarkly highlights regressions and pauses a guarded rollout automatically if needed.


Monitor an AI Config

1. Navigate to the detail page for the AI Config. Select the **Monitoring** tab.
2. At the top of the tab, use the controls to specify the monitoring data you want to view:
 - Use the **Environment** dropdown to select the [environment](#) you want to monitor. Performance metrics are specific to each environment.
 - Select a date range.
 - Use the **Charts** dropdown to select which set of charts to view.
3. Review charts of the available monitoring data. The data in each chart applies to the selected environment and is split out by variation:
 - The **Tokens** charts show include average input and output tokens used by the AI Config in this environment, by variation.
 - The **Satisfaction** chart is the percentage of “thumbs up” ratings provided by end users who have encountered the AI Config in this environment, by variation.
 - The **Generations** is the average number of successful generations completed using the AI Config in this environment, by variation.
 - The **Time to generate** is the average duration per generation. In other words, this is the total duration of calls to your LLM provider divided by the number of total generations using the AI Config in this environment, by variation.
 - **Error rate** is the percentage of errors out of the total number of generations attempted for the AI Config in this environment, by variation.
 - **Time to first token** is the mean time it takes to generate the initial token.
 - **Costs** charts show the sum of the input token cost and output tokens cost used by the AI Config in this environment, by variation.
4. Review the table of available monitoring data. The table includes the data displayed in the charts as well as additional data. For example, the table includes both the average and the total input and output tokens.
5. (Optional) In the table, select the variations or versions for which you want to view data. The charts update based on your selection.
6. (Optional) Click **Export data as CSV** to download a CSV file for further analysis.

[Tracking AI metrics.](#)

The data on the **Monitoring** tab updates approximately each minute. For each metric, the results are broken out by AI Config variation. Metrics with no results display “No data.” If there is no data for a particular variation, that variation is not included in the total displayed at the top of the metric card.

Here is a partial image of the **Monitoring** tab:



Was this page helpful?

☐ Yes

☐ No

[< Previous](#)

[Manage AI Configs](#)

[Next >](#)

© 2025 Catamorphic Co.